

Network Programming for Windows 05: **Winsock I/O Methods**

jintaeks@dongseo.ac.kr

Division of Digital Contents, DSU

May 2019

Outline

- ✓ Socket Modes
- ✓ Socket I/O Models
- ✓ I/O Model Consideration

-
- ✓ Winsock features socket modes and socket I/O models to control how I/O is processed on a socket.
 - ✓ A **socket mode** simply determines how Winsock functions behave when called with a socket.
 - Two socket modes: **blocking** and **non-blocking**.
 - ✓ A **socket model** describes how an application manages and processes I/O on a socket.
 - **blocking**, **select()**, **WSAAsyncSelect()**, **WSAEventSelect()**, **overlapped I/O**, and **completion port**.

Socket Modes

- ✓ In **blocking mode**, Winsock calls that perform I/O, such as send() and recv() wait until the operation is complete before they return to the program.
- ✓ In **non-blocking mode**, the Winsock functions return immediately.

Blocking Mode

- ✓ The problem with this code is that the `recv()` function might never return.

```
SOCKET sock;
char buff[256];
int done = 0, nBytes;
...

while (!done)
{
    nBytes = recv(sock, buff, 65);
    if (nBytes == SOCKET_ERROR)
    {
        printf("recv failed with error %d\n", WSAGetLastError());
        return;
    }
    DoComputationOnData(buff);
}
```

- ✓ One method is to **separate the application into a reading thread and a computation thread.**

```
#define MAX_BUFFER_SIZE    4096

// Initialize critical section (data) and create
// an auto-reset event (hEvent) before creating the two threads
CRITICAL_SECTION data;
HANDLE          hEvent;

SOCKET          sock;
TCHAR          buff[MAX_BUFFER_SIZE];
int              done = 0;

// Create and connect sock
...
```



```
// Reader thread
void ReadThread(void)
{
    int nTotal = 0,
        nRead = 0,
        nLeft = 0,
        nBytes = 0;

    while (!done)    {
        nTotal = 0;
        nLeft = NUM_BYTES_REQUIRED;

        // However many bytes constitutes enough data for processing (i.e. non-zero)
        while (nTotal != NUM_BYTES_REQUIRED)    {
            EnterCriticalSection(&data);
            nRead = recv(sock, &(buff[MAX_BUFFER_SIZE - nBytes]), nLeft, 0);
            if (nRead == -1)    {
                printf("error\n");
                ExitThread();
            }
            nTotal += nRead;
            nLeft -= nRead;

            nBytes += nRead;
            LeaveCriticalSection(&data);
        }
        SetEvent(hEvent);
    }
}
```

```
// Computation thread
void ProcessThread(void)
{
    WaitForSingleObject(hEvent);
    EnterCriticalSection(&data);
    DoSomeComputationOnData(buff);

    // Remove the processed data from the input
    // buffer, and shift the remaining data to the start of the array
    nBytes -= NUM_BYTES_REQUIRED;

    LeaveCriticalSection(&data);
}
```


Non-blocking Mode

- ✓ Winsock API calls that deal with sending and receiving data or connection management return immediately.

```
SOCKET      s;  
unsigned long ul = 1;  
int          nRet;  
  
s = socket(AF_INET, SOCK_STREAM, 0);  
  
nRet = ioctlsocket(s, FIONBIO, (unsigned long *)&ul);  
  
if (nRet == SOCKET_ERROR) {  
    // Failed to put the socket into non-blocking mode  
}
```

- ✓ Because non-blocking calls frequently fail with the **WSAEWOULDBLOCK** error, you should check all return codes and be prepared for failure at any time.

Socket I/O Models

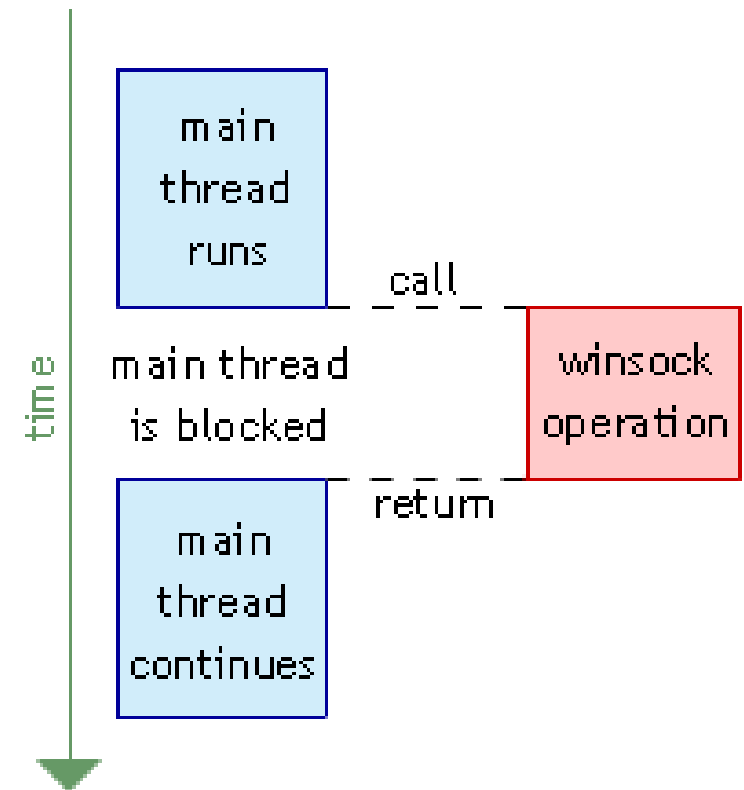
A **socket model** describes how an application manages and processes I/O on a socket.

- ① blocking
- ② select()
- ③ WSAAsyncSelect()
- ④ WSAEventSelect()
- ⑤ overlapped
- ⑥ completion port

Model	Blocking mode	Notification method		
		<i>none</i>	<i>on network event</i>	<i>on completion</i>
Blocking sockets	blocking	x		
Polling	non-blocking	x		
Select	both		blocking select	
WSAAsyncSelect	non-blocking		window message	
WSAEventSelect	non-blocking		event objects	
Overlapped I/O: blocking	N/A			blocking call
Overlapped I/O: polling	N/A	x		
Overlapped I/O: completion routines	N/A			callback function
Overlapped I/O: completion ports	N/A			completion port

The Blocking Model

- ✓ Straightforward model.
 - chapter 1
- ✓ Applications following this model typically use one or two **threads per socket connection** for handling I/O.
- ✓ Each thread will then issue blocking operations, such as `send()` and `recv()`.
- ✓ For very simple applications and **rapid prototyping**, this model is very useful.



```

DWORD WINAPI ServerListenThread(LPVOID lpParam)
{
    CONNECTION_OBJ *ConnObj = NULL;
    HANDLE          hThread = NULL;
    SOCKET          s;
    int             rc;

    s = (SOCKET)lpParam;
    ...
    while (1)
    {
        // Wait for an incoming client connection
        ns = accept(s, (SOCKADDR *)&saAccept, &acceptlen);
        if (ns == INVALID_SOCKET)
        {
            fprintf(stderr, "accept failed: %d\n", WSAGetLastError());
            return -1;
        }

        InterlockedIncrement(&gConnectedClients);
    }
}

```

...

```
// Allocate a connection object for this client
```

```
ConnObj = GetConnectionObj(ns);
```

```
// Create a receiver thread for this connection
```

```
hThread = CreateThread(NULL, 0, ReceiveThread, (LPVOID)ConnObj, 0, NULL);
```

```
if (hThread == NULL)
```

```
{
```

```
    fprintf(stderr, "CreateThread failed: %d\n", GetLastError());
```

```
    ExitThread(-1);
```

```
}
```

```
CloseHandle(hThread);
```

```
// Create a sender thread for this connection
```

```
hThread = CreateThread(NULL, 0, SendThread, (LPVOID)ConnObj, 0, NULL);
```

```
if (hThread == NULL)
```

```
{
```

```
    fprintf(stderr, "CreateThread failed: %d\n", GetLastError());
```

```
    ExitThread(-1);
```

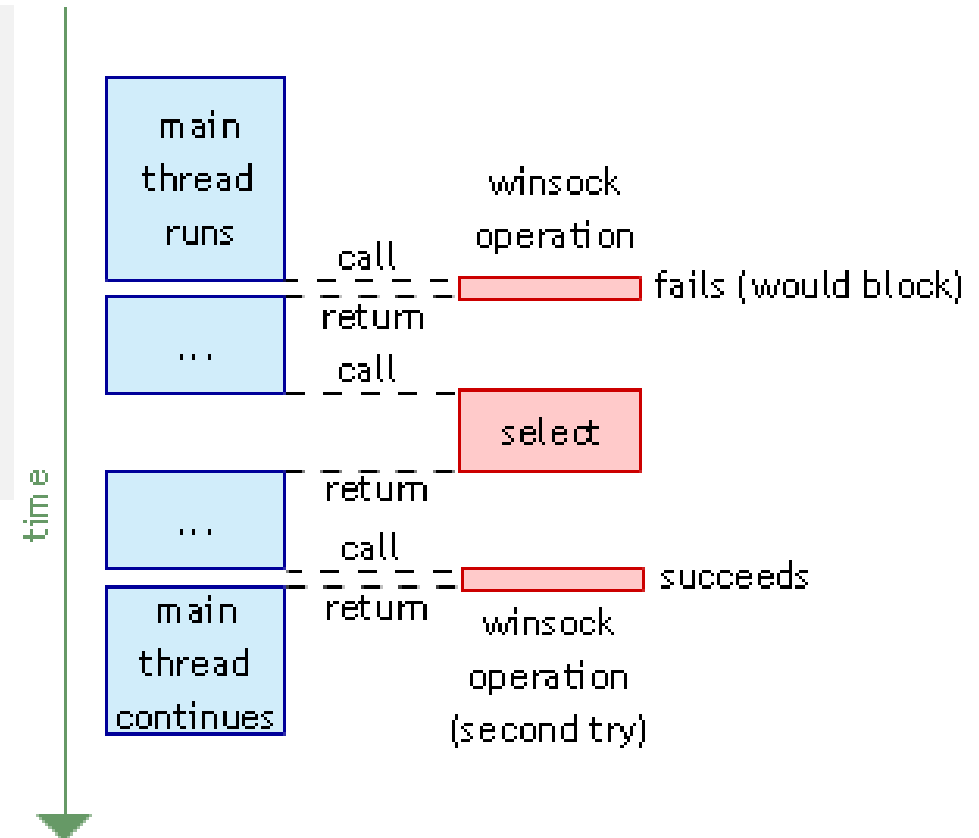
```
}
```

```
CloseHandle(hThread);
```

The select Model

- ✓ We call it the select model because it centers on using the **select()** function to manage I/O.
- ✓ The **select()** function can be used to determine if there is data on a socket and if a socket can be written to.

```
int select(  
    int nfd,  
    fd_set FAR * readfds,  
    fd_set FAR * writefds,  
    fd_set FAR * exceptfds,  
    const struct timeval FAR * timeout  
);
```



```

SOCKET s;
fd_set fdread;
int ret;

// Create a socket, and accept a connection

// Manage I/O on the socket
while (TRUE){
    // Always clear the read set before calling select()
    FD_ZERO(&fdread);
    // Add socket s to the read set
    FD_SET(s, &fdread);

    if ((ret = select(0, &fdread, NULL, NULL, NULL)) == SOCKET_ERROR) {
        // Error condition
    }

    if (ret > 0) {
        // For this simple case, select() should return the value 1.
        // An application dealing with more than one socket could get a value
        // greater than 1. At this point, your
        // application should check to see whether the socket is part of a set.
        if (FD_ISSET(s, &fdread)) {
            // A read event has occurred on socket s
        }
    }
}

```


Practice

- ✓ chapter05 → select client/server projects
- ✓ chapter05 → blocking server project

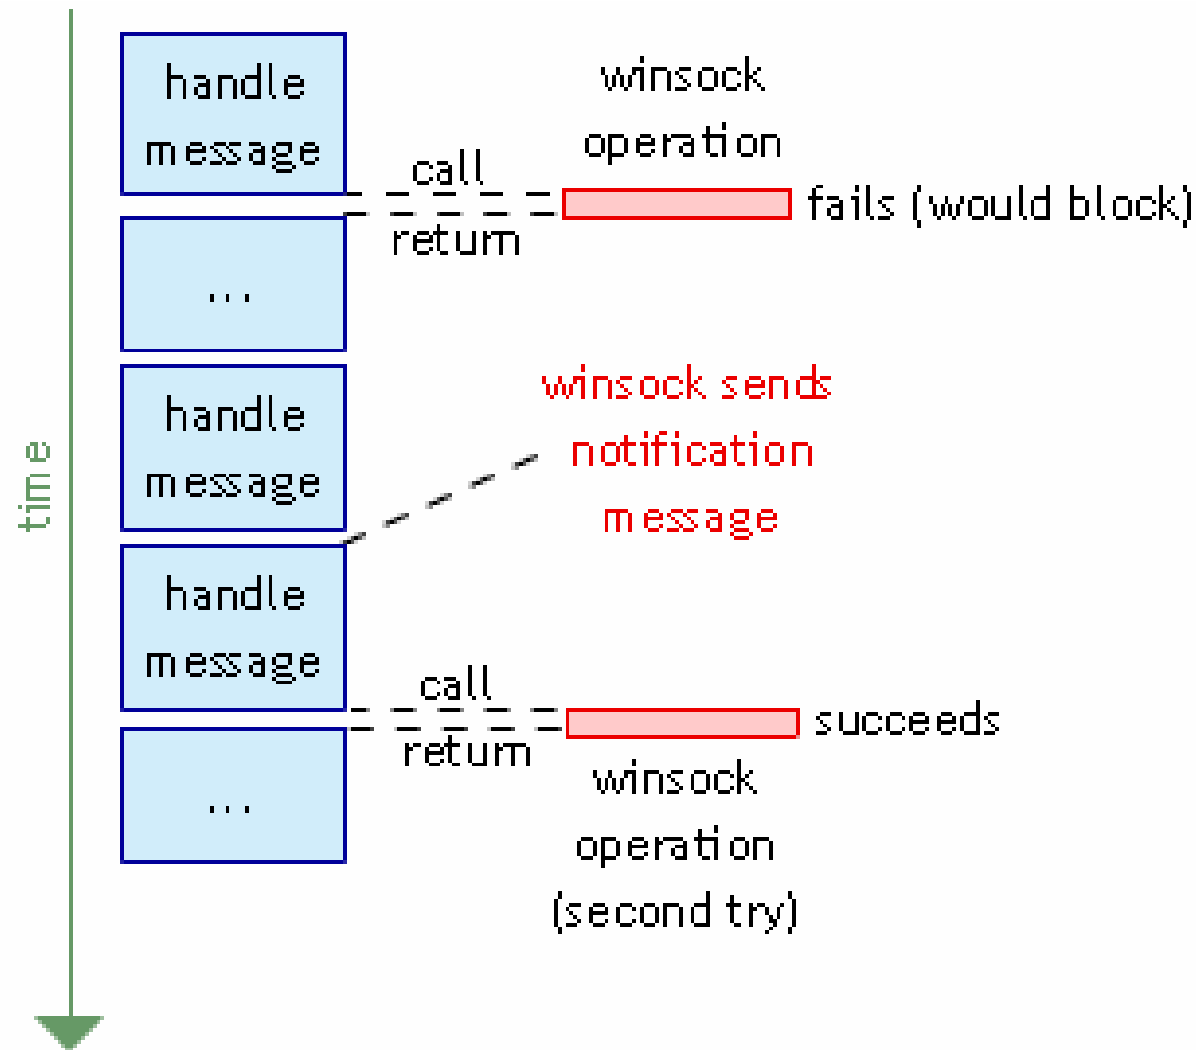
The WSAAsyncSelect() Model

- ✓ Asynchronous I/O model that allows an application to receive **Windows message-based** notification of network events on a socket.
- ✓ The WSAAsyncSelect() and WSAEventSelect() models **does not** provide asynchronous data transfer like the overlapped and completion port models.
- ✓ This model is also used by the Microsoft Foundation Class (**MFC**) **CSocket** object.

```
int WSAAsyncSelect(  
    SOCKET s,  
    HWND hWnd,  
    unsigned int wMsg,  
    long lEvent  
);
```

Practice

✓ chapter05 → AsyncSelectServer project



case 1




1024

time



case *FD_READ*:


WSARecv(, 1024, ...)



1024


case *FD_WRITE*:

WSASend(, 1024, ...)



1024


case 2



time


case *FD_READ*:

WSARecv(, 1024, ...)




case *FD_WRITE*:

WSASend(, 512, ...)




case *FD_WRITE*:

WSASend(, 512, ...)




case 3



time

case *FD_READ*:


WSARecv(, 512, ...)



512

case *FD_WRITE*:


WSASend(, 512, ...)



512

case *FD_READ*:


WSARecv(, 512, ...)



512


case *FD_WRITE*:

WSASend(, 512, ...)




512

case 4




time


case *FD_READ*:
WSARecv(, 1024, ...)



case *FD_WRITE*:
WSASend(, 512, ...)




case *FD_READ*:
SocketInfo->RecvPosted = *TRUE*;
return 0;
case *FD_WRITE*:




```
case FD_READ:
    SocketInfo->RecvPosted = TRUE;
    return 0;
```

```
case FD_WRITE:
    WSASend(, 512, ...)
```



```
PostMessage(hwnd, WM_SOCKET,
    wParam, FD_READ);
```

```
case FD_READ:
    WSARecv(, 1024, ...)
```




```
LRESULT CALLBACK WindowProc(HWND hwnd, UINT uMsg, WPARAM wParam, LPARAM lParam)
```

```
{
```

```
    SOCKET Accept;
```

```
    LPSOCKET_INFORMATION SocketInfo;
```

```
    DWORD RecvBytes;
```

```
    DWORD SendBytes;
```

```
    DWORD Flags;
```

```
    if (uMsg == WM_SOCKET)
```

```
    {
```

```
        if (WSAGETSELECTERROR(lParam))
```

```
        {
```

```
            printf("Socket failed with error %d\n", WSAGETSELECTERROR(lParam));
```

```
            FreeSocketInformation(wParam);
```

```
        }
```

```
    else
```

```
    {
```

```
        printf("Socket looks fine!\n");
```

```
        switch (WSAGETSECTEVENT(lParam))
```

```
        {
```

```
            case FD_ACCEPT:
```

```
                if ((Accept = accept(wParam, NULL, NULL)) == INVALID_SOCKET)
```

```
                {
```

```
                    printf("accept() failed with error %d\n", WSAGetLastError());
```

```
                    break;
```

```
                }
```

```
            else
```

```
                printf("accept() is OK!\n");
```

```
            // Create a socket information structure to associate with the socket for processing
```

```
            CreateSocketInformation(Accept);
```

```
            printf("Socket number %d connected\n", Accept);
```

```
            WSAAsyncSelect(Accept, hwnd, WM_SOCKET, FD_READ | FD_WRITE | FD_CLOSE);
```

```
            break;
```

I/O

```
case FD_READ:
```

```
SocketInfo = GetSocketInformation(wParam);  
// Read data only if the receive buffer is empty  
if (SocketInfo->BytesRECV != 0)  
{  
    SocketInfo->RecvPosted = TRUE;  
    return 0;  
}
```

```
else
```

```
{  
    SocketInfo->DataBuf.buf = SocketInfo->Buffer;  
    SocketInfo->DataBuf.len = DATA_BUFSIZE;
```

```
Flags = 0;
```

```
if (WSARecv(SocketInfo->Socket, &(SocketInfo->DataBuf), 1, &RecvBytes,  
    &Flags, NULL, NULL) == SOCKET_ERROR)
```

```
{
```

```
    if (WSAGetLastError() != WSAEWOULDBLOCK)
```

```
    {
```

```
        printf("WSARecv() failed with error %d\n", WSAGetLastError());
```

```
        FreeSocketInformation(wParam);
```

```
        return 0;
```

```
    }
```

```
}
```

```
else // No error so update the byte count
```

```
{
```

```
    printf("WSARecv() is OK!\n");
```

```
    SocketInfo->BytesRECV = RecvBytes;
```

```
}
```

```
}
```

```
// DO NOT BREAK HERE SINCE WE GOT A SUCCESSFUL RECV. Go ahead
```

```
// and begin writing data to the client
```

```
case FD_WRITE:
```

```
SocketInfo = GetSocketInformation(wParam);
```

```
case FD_WRITE:
```

```
SocketInfo = GetSocketInformation(wParam);
```

```
if (SocketInfo->BytesRECV > SocketInfo->BytesSEND)
```

```
{
```

```
    SocketInfo->DataBuf.buf = SocketInfo->Buffer + SocketInfo->BytesSEND;
```

```
    SocketInfo->DataBuf.len = SocketInfo->BytesRECV - SocketInfo->BytesSEND;
```

```
    if (WSASend(SocketInfo->Socket, &(SocketInfo->DataBuf), 1, &SendBytes, 0,  
        NULL, NULL) == SOCKET_ERROR)
```

```
    {
```

```
        if (WSAGetLastError() != WSAEWOULDBLOCK)
```

```
        {
```

```
            printf("WSASend() failed with error %d\n", WSAGetLastError());
```

```
            FreeSocketInformation(wParam);
```

```
            return 0;
```

```
        }
```

```
    }
```

```
else // No error so update the byte count
```

```
{
```

```
    printf("WSASend() is OK!\n");
```

```
    SocketInfo->BytesSEND += SendBytes;
```

```
}
```

```
}
```

```
if (SocketInfo->BytesSEND == SocketInfo->BytesRECV)
```

```
{
```

```
    SocketInfo->BytesSEND = 0;
```

```
    SocketInfo->BytesRECV = 0;
```

```
    // If a RECV occurred during our SENDs then we need to post an FD_READ
```

```
notification on the socket
```

```
    if (SocketInfo->RecvPosted == TRUE)
```

```
    {
```

```
        SocketInfo->RecvPosted = FALSE;
```

```
        PostMessage(hwnd, WM_SOCKET, wParam, FD_READ);
```

The ~~WSAEventSelect~~ Model

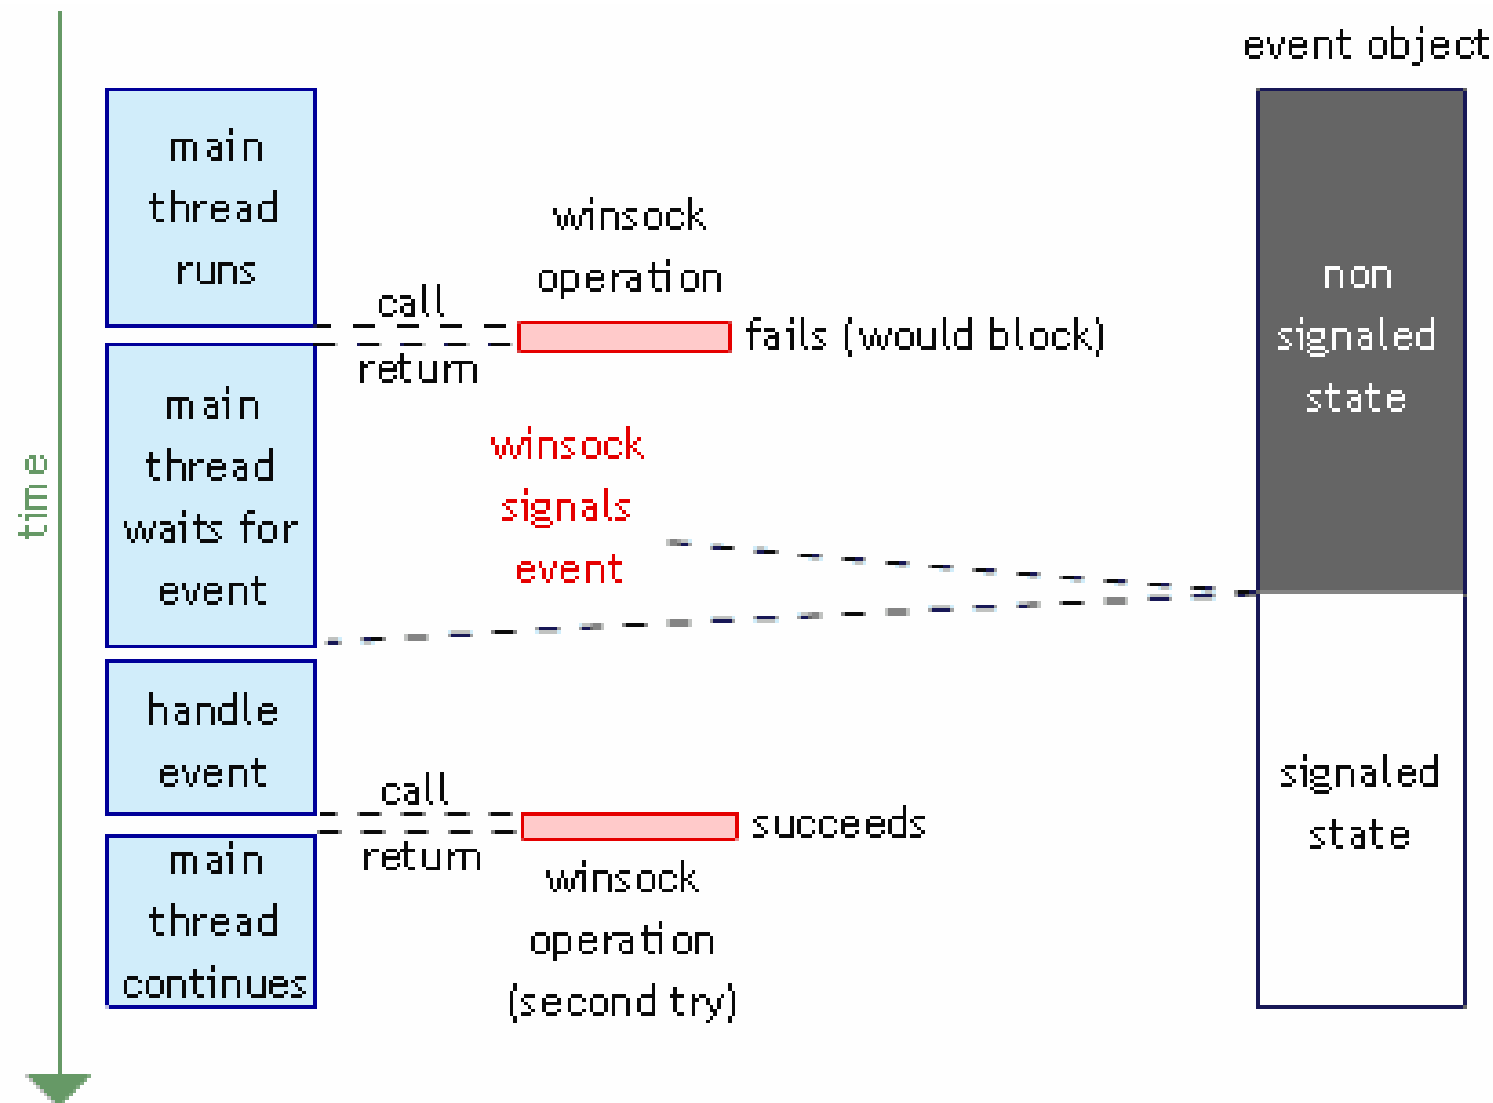
- ✓ Network events are notified via an **event object** handle instead of a window procedure.

```
int WSAEventSelect(  
    SOCKET s,  
    WSAEVENT hEventObject,  
    long lNetworkEvents  
);
```

```
DWORD WSAWaitForMultipleEvents(  
    DWORD cEvents,  
    const WSAEVENT FAR * lphEvents,  
    BOOL fWaitAll,  
    DWORD dwTimeout,  
    BOOL fAlertable  
);
```

Practice

✓ chapter05 → EventSelectServer project



The Overlapped Model

- ✓ The overlapped model's basic design allows your application to post one or more asynchronous I/O requests at a time using an overlapped data structure.
- ✓ The model's overall design is based on the **Windows overlapped I/O** mechanisms.
- ✓ To use the overlapped I/O model on a socket, you must first **create** a socket that has the **overlapped flag** set.
 - `if ((ListenSocket = WSASocket(AF_INET, SOCK_STREAM, 0, NULL, 0, WSA_FLAG_OVERLAPPED)) == INVALID_SOCKET)`
- ✓ To use overlapped I/O, each WSA function takes a **WSAOVERLAPPED** structure as a parameter.
 - `if (WSARecv(SI->Socket, &(SI->DataBuf), 1, &RecvBytes, &Flags, &(SI->Overlapped), NULL) == SOCKET_ERROR)`

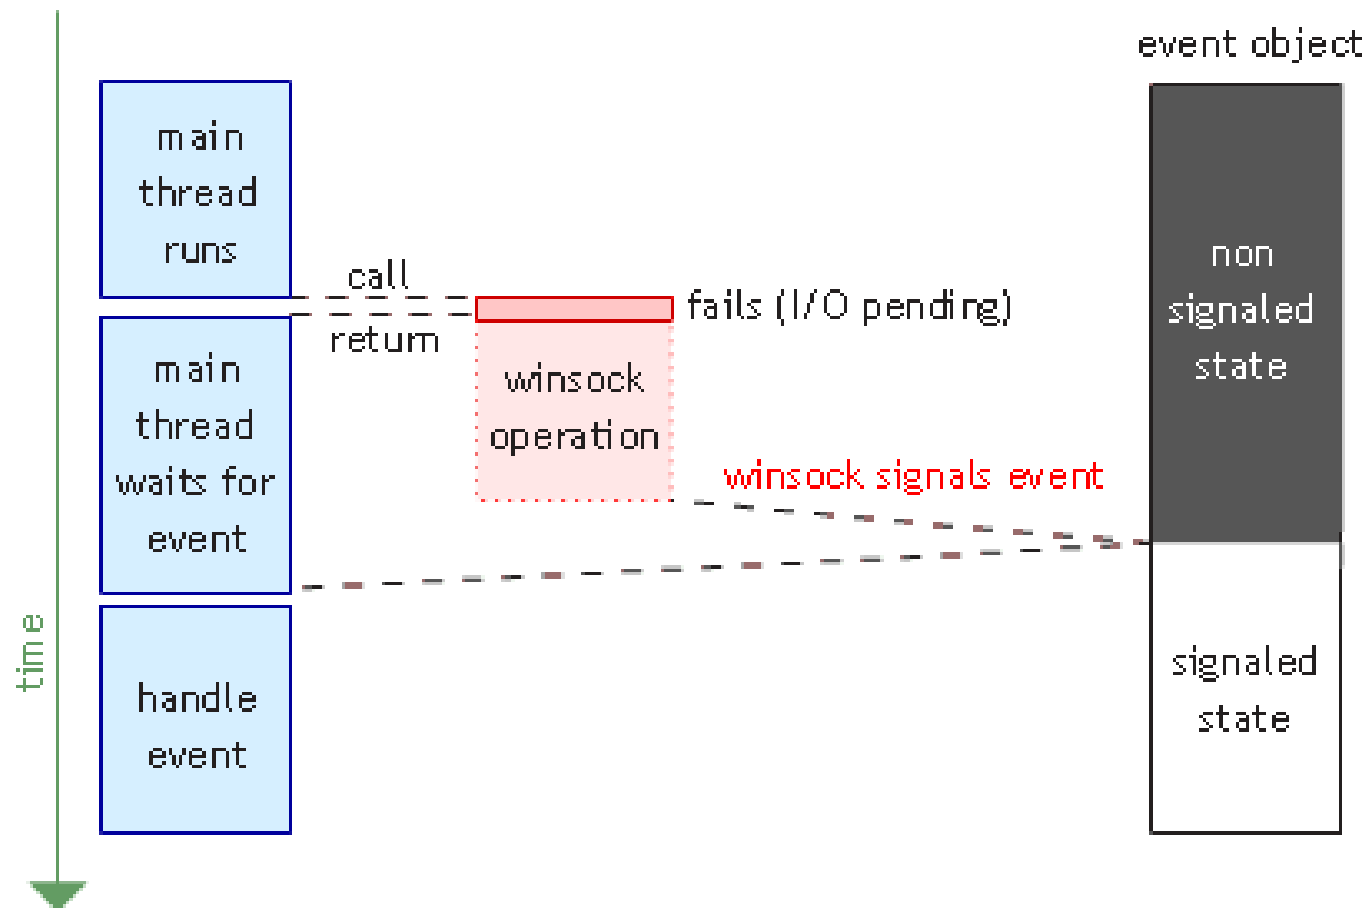
Event Notification

- ✓ The event notification method of overlapped I/O requires associating **Windows event objects** with **WSAOVERLAPPED** structures.
- ✓ When I/O calls such as WSARecv() and WSASend() are made using a WSAOVERLAPPED structure, they return immediately.

```
typedef struct WSAOVERLAPPED
{
    DWORD    Internal;
    DWORD    InternalHigh;
    DWORD    Offset;
    DWORD    OffsetHigh;
    WSAEVENT hEvent;
} WSAOVERLAPPED, FAR * LPWSAOVERLAPPED;
```

```
BOOL WSAGetOverlappedResult(  
    SOCKET s,  
    LPWSAOVERLAPPED lpOverlapped,  
    LPDWORD lpcbTransfer,  
    BOOL fWait,  
    LPDWORD lpdwFlags  
);
```

- ✓ If the **WSAGetOverlappedResult()** function succeeds, the return value is TRUE.
- ✓ If the return value is FALSE, one of the following statements is true:
 - The overlapped I/O operation is still pending.
 - The overlapped operation completed, but with errors.



1. **Create** a socket and begin **listening** for a connection on a specified port.
2. **Accept** an inbound connection.
3. Create a **WSAOVERLAPPED structure** for the accepted socket and assign an **event object** handle to the structure. Also assign the **event object handle** to an event array to be used later by the `WSAWaitForMultipleEvents()` function.
4. Post an asynchronous **WSARecv()** request on the socket by specifying the WSAOVERLAPPED structure as a parameter.

5. Call **WSAWaitForMultipleEvents()** using the **event array** and wait for the event associated with the overlapped call to become signaled.
6. Determine the **return status** of the overlapped call by using **WSAGetOverlappedResult()**.
7. **Reset the event** object by using **WSAResetEvent()** with the event array and process the completed overlapped request.
8. **Post another overlapped WSARecv()** request on the socket.
9. Repeat steps 5–8.

→ *DWORD WINAPI* ProcessIO(*LPVOID lpParameter*)

```

#define PORT 5150
#define DATA_BUFSIZE 8192

typedef struct _SOCKET_INFORMATION {
    CHAR Buffer[DATA_BUFSIZE];
    WSABUF DataBuf;
    SOCKET Socket;
    WSAOVERLAPPED Overlapped;
    DWORD BytesSEND;
    DWORD BytesRECV;
} SOCKET_INFORMATION, *LPSOCKET_INFORMATION;

DWORD WINAPI ProcessIO(LPVOID lpParameter);

DWORD EventTotal = 0;
WSAEVENT EventArray[WSA_MAXIMUM_WAIT_EVENTS];
LPSOCKET_INFORMATION SocketArray[WSA_MAXIMUM_WAIT_EVENTS];
CRITICAL_SECTION CriticalSection;

int main(int argc, char **argv)
{
    WSADATA wsaData;
    SOCKET ListenSocket, AcceptSocket;
    SOCKADDR_IN InternetAddr;

```

```

int main(int argc, char **argv)
{
    WSADATA wsaData;
    SOCKET ListenSocket, AcceptSocket;
    SOCKADDR_IN InternetAddr;
    DWORD Flags;
    DWORD ThreadId;
    DWORD RecvBytes;

    InitializeCriticalSection(&CriticalSection);

```

```

WSAStartup((2, 2), &wsaData);
printf("WSAStartup() looks nice!\n");

```

Step1

```

ListenSocket = WSASocket(AF_INET, SOCK_STREAM, 0, NULL, 0, WSA_FLAG_OVERLAPPED);
printf("WSASocket() is OK lol!\n");

```

```

InternetAddr.sin_family = AF_INET;
InternetAddr.sin_addr.s_addr = htonl(INADDR_ANY);
InternetAddr.sin_port = htons(PORT);

```

```

bind(ListenSocket, (PSOCKADDR)&InternetAddr, sizeof(InternetAddr));
printf("YOu see, bind() is working!\n");

```

```

listen(ListenSocket, 5);
printf("listen() is OK maa...\n");

```

```

EventArray[0] = WSACreateEvent();
printf("WSACreateEvent() is OK!\n");

```

```

// Create a thread to service overlapped requests
CreateThread(NULL, 0, ProcessIO, NULL, 0, &ThreadId);
printf("Nothing to say, CreateThread() is OK!\n");

```

```

EventTotal = 1;

```

```
EventTotal = 1;
```

```
while (TRUE)
```

```
{
```

```
// Accept inbound connections
```

```
AcceptSocket = accept(ListenSocket, NULL, NULL);
```

```
printf("accept() is OK!\n");
```

Step2

```
EnterCriticalSection(&CriticalSection);
```

Step3

```
// Create a socket information structure to associate with the accepted socket  
SocketArray[EventTotal] = (LPSOCKET_INFORMATION)GlobalAlloc(GPTR,  
sizeof(SOCKET_INFORMATION));
```

```
printf("GlobalAlloc() for LPSOCKET_INFORMATION is pretty fine!\n");
```

```
// Fill in the details of our accepted socket
```

```
SocketArray[EventTotal]->Socket = AcceptSocket;
```

```
ZeroMemory(&(SocketArray[EventTotal]->Overlapped), sizeof(OVERLAPPED));
```

```
SocketArray[EventTotal]->BytesSEND = 0;
```

```
SocketArray[EventTotal]->BytesRECV = 0;
```

```
SocketArray[EventTotal]->DataBuf.Len = DATA_BUFSIZE;
```

```
SocketArray[EventTotal]->DataBuf.buf = SocketArray[EventTotal]->Buffer;
```

```
SocketArray[EventTotal]->Overlapped.hEvent = EventArray[EventTotal] = WSACreateEvent();
```

```
printf("WSACreateEvent() is OK!\n");
```

```
// Post a WSAREcv() request to to begin receiving data on the socket
```

```
Flags = 0;
```

```
WSAREcv(SocketArray[EventTotal]->Socket,
```

```
&(SocketArray[EventTotal]->DataBuf), 1, &RecvBytes, &Flags, &(SocketArray[EventTotal]-  
>Overlapped), NULL);
```

```
printf("WSAREcv() should be working!\n");
```

```

// Post a WSARecv() request to to begin receiving data on the socket
Flags = 0;
WSARecv(SocketArray[EventTotal]->Socket,
        &(SocketArray[EventTotal]->DataBuf), 1, &RecvBytes, &Flags, &(SocketArray[EventTotal]-
>Overlapped), NULL);
printf("WSARecv() should be working!\n");

EventTotal++;
LeaveCriticalSection(&CriticalSection);

// Signal the first event in the event array to tell the worker thread to
// service an additional event in the event array
WSASetEvent(EventArray[0]);
printf("Don't worry, WSASetEvent() is OK!\n");
}
}

```

Step4

```
DWORD WINAPI ProcessIO(LPVOID lpParameter)
```

```
{
```

```
    DWORD Index;
```

```
    DWORD Flags;
```

```
    LPSOCKET_INFORMATION SI;
```

```
    DWORD BytesTransferred;
```

```
    DWORD i;
```

```
    DWORD RecvBytes, SendBytes;
```

```
    // Process asynchronous WSASend, WSARecv requests
```

```
    while (TRUE)
```

```
    {
```

```
        Index = WSAWaitForMultipleEvents(EventTotal, EventArray, FALSE, WSA_INFINITE, FALSE);
```

```
        printf("WSAWaitForMultipleEvents() is OK!\n");
```

```
        // If the event triggered was zero then a connection attempt was made
```

```
        // on our listening socket.
```

```
        if ((Index - WSA_WAIT_EVENT_0) == 0)
```

```
        {
```

```
            WSAResetEvent(EventArray[0]);
```

```
            continue;
```

```
        }
```

```
        SI = SocketArray[Index - WSA_WAIT_EVENT_0];
```

```
        WSAResetEvent(EventArray[Index - WSA_WAIT_EVENT_0]);
```

```
        WSAGetOverlappedResult(SI->Socket, &(SI->Overlapped), &BytesTransferred, FALSE, &Flags);
```

```
        if( BytesTransferred == 0)
```

```
        {
```

```
            printf("Closing socket %d\n", SI->Socket);
```

```
            closesocket(SI->Socket);
```

```
            printf("closesocket() is OK!\n");
```

Step5

Step6,7


```

printf("closesocket() is OK!\n");

GlobalFree(SI);
WSACloseEvent(EventArray[Index - WSA_WAIT_EVENT_0]);
// Cleanup SocketArray and EventArray by removing the socket event handle
// and socket information structure if they are not at the end of the arrays
EnterCriticalSection(&CriticalSection);

if ((Index - WSA_WAIT_EVENT_0) + 1 != EventTotal)
    for (i = Index - WSA_WAIT_EVENT_0; i < EventTotal; i++)
    {
        EventArray[i] = EventArray[i + 1];
        SocketArray[i] = SocketArray[i + 1];
    }

EventTotal--;
LeaveCriticalSection(&CriticalSection);
continue;
}
// Check to see if the BytesRECV field equals zero. If this is so, then
// this means a WSARcv call just completed so update the BytesRECV field
// with the BytesTransferred value from the completed WSARcv() call.
if (SI->BytesRECV == 0)
{
    SI->BytesRECV = BytesTransferred;
    SI->BytesSEND = 0;
}
else
{
    SI->BytesSEND += BytesTransferred;
}

```

```

if (SI->BytesSEND < SI->BytesRECV)
{
    // Post another WSASend() request.
    // Since WSASend() is not guaranteed to send all of the bytes requested,
    // continue posting WSASend() calls until all received bytes are sent
    ZeroMemory(&(SI->Overlapped), sizeof(WSAOVERLAPPED));
    SI->Overlapped.hEvent = EventArray[Index - WSA_WAIT_EVENT_0];

    SI->DataBuf.buf = SI->Buffer + SI->BytesSEND;
    SI->DataBuf.len = SI->BytesRECV - SI->BytesSEND;

    WSASend(SI->Socket, &(SI->DataBuf), 1, &SendBytes, 0, &(SI->Overlapped), NULL);
    printf("WSASend() is OK!\n");
}
else
{
    SI->BytesRECV = 0;
    // Now that there are no more bytes to send post another WSARecv() request
    Flags = 0;
    ZeroMemory(&(SI->Overlapped), sizeof(WSAOVERLAPPED));
    SI->Overlapped.hEvent = EventArray[Index - WSA_WAIT_EVENT_0];

    SI->DataBuf.len = DATA_BUFSIZE;
    SI->DataBuf.buf = SI->Buffer;

    WSARecv(SI->Socket, &(SI->DataBuf), 1, &RecvBytes, &Flags, &(SI->Overlapped), NULL);
    printf("WSARecv() is OK!\n");
}
}
}

```

Step8

```

WSARecv(SI->Socket, &(SI->DataBuf), 1, &RecvBytes, &Flags, &(SI->Overlapped), NULL);
printf("WSARecv() is OK!\n");

```

Question: What's this code?

```
if (listen(ListenSocket, 5))
{
    printf("listen() failed with error %d\n", WSAGetLastError());
    return 1;
}
else
    printf("listen() is OK maa...\n");

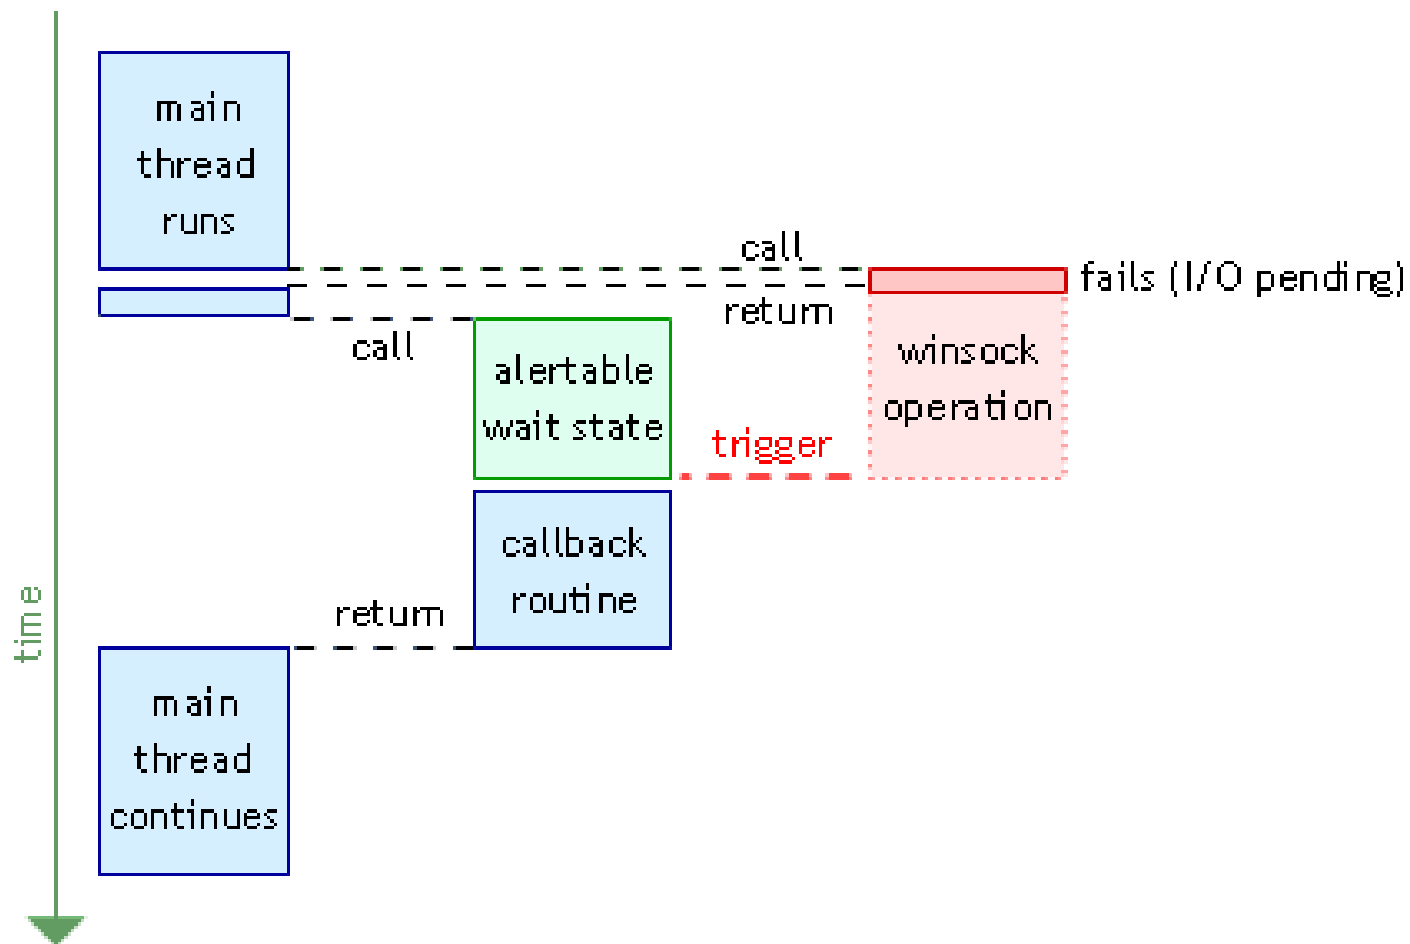
// Setup the listening socket for connections
if ((AcceptSocket = WSASocket(AF_INET, SOCK_STREAM, 0, NULL, 0, WSA_FLAG_OVERLAPPED)) ==
INVALID_SOCKET)
{
    printf("Failed to get a socket %d\n", WSAGetLastError());
    return 1;
}
else
    printf("WSASocket() looks OK!\n");

if ((EventArray[0] = WSACreateEvent()) == WSA_INVALID_EVENT)
{
    printf("WSACreateEvent() failed with error %d\n", WSAGetLastError());
}
```

Completion Routines

- ✓ **Completion routines** are the other method your application can use to **manage completed overlapped I/O requests**.
- ✓ Completion routines are simply **functions** that the system invokes when an overlapped I/O request completes.
- ✓ Their primary role is to service a completed I/O request using the caller's thread. In addition, applications can continue overlapped I/O processing through the completion routine.
- ✓ To use completion routines for overlapped I/O requests, your application must specify a completion routine, along with a **WSAOVERLAPPED** structure.

```
void CALLBACK CompletionROUTINE(  
    DWORD dwError,  
    DWORD cbTransferred,  
    LPWSAOVERLAPPED lpOverlapped,  
    DWORD dwFlags  
);
```



1. **Create a socket** and begin **listening** for a connection on a specified port.
2. **Accept** an inbound connection.
3. Create a **WSAOVERLAPPED** structure for the accepted socket.
4. Post an asynchronous **WSARecv** request on the socket by ① specifying the WSAOVERLAPPED structure as a parameter and ② **supplying a completion routine.**

5. Call **WSAWaitForMultipleEvents()** with the fAlertable parameter set to TRUE and wait for an overlapped request to complete.

When an overlapped request completes, the **completion routine automatically executes** and WSAWaitForMultipleEvents() returns WSA_IO_COMPLETION. Inside the completion routine, then post another overlapped WSARecv request with a completion routine.

6. Verify that **WSAWaitForMultipleEvents()** returns **WSA_IO_COMPLETION**.


7. Repeat steps 5 and 6.

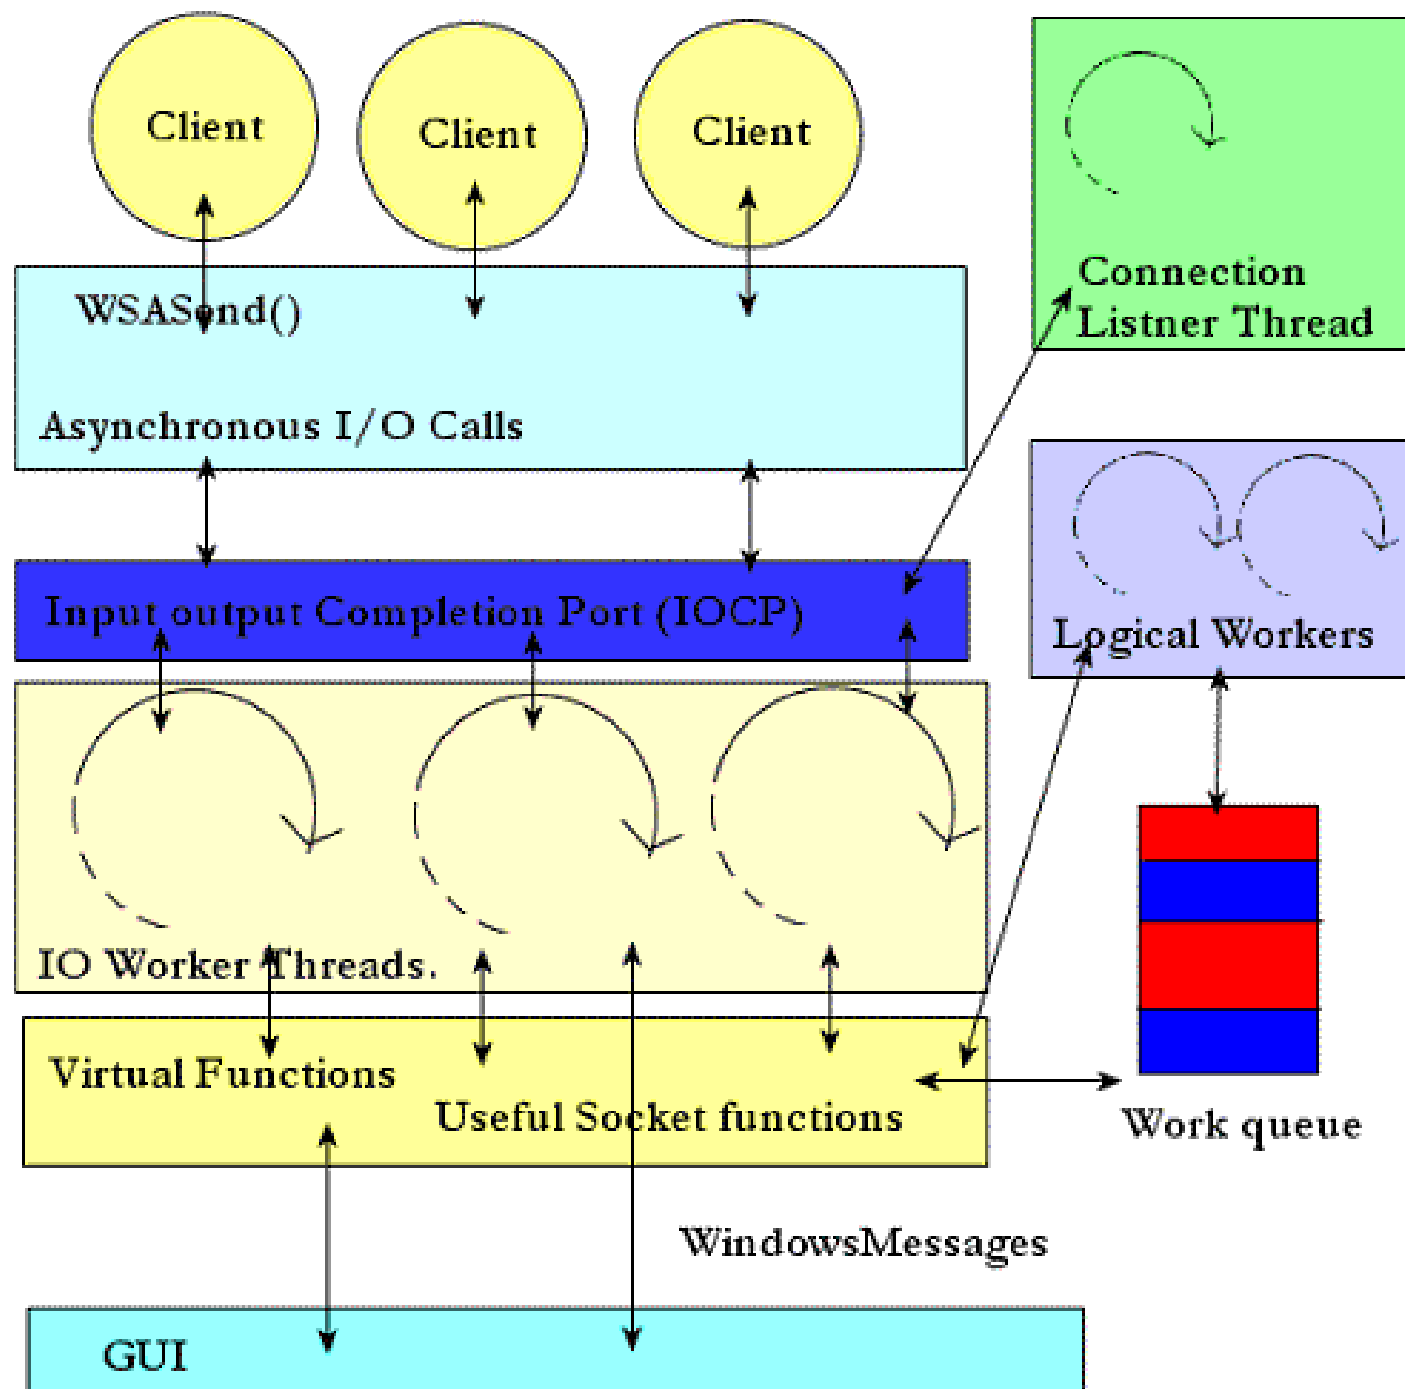
IOCP Model

FINALLY! The Completion Port Model

- ✓ The **completion port model** offers the best system performance possible when an application has to manage many sockets at once.
- ✓ It's available Windows NT and above.
- ✓ To begin using this model, you are required to create an **I/O completion port object** that will be used to manage multiple I/O requests for any number of socket handles.

```
HANDLE CreateIoCompletionPort(  
    HANDLE FileHandle,  
    HANDLE ExistingCompletionPort,  
    DWORD CompletionKey,  
    DWORD NumberOfConcurrentThreads  
);
```





Worker Threads and Completion Ports

- ✓ After a **completion port** is successfully **created**, you can begin to **associate socket handles** with the object.
- ✓ Before associating sockets, though, you have to create **one or more worker threads** to service the completion port when socket I/O requests are posted to the completion port object.
- ✓ Calling the **CreateIoCompletionPort()** function on an existing completion port and supplying the first three parameters with socket information.

1. Create a completion port.

1. The fourth parameter is left as 0, specifying that only one worker thread per processor will be allowed.
2. Determine how many **processors** exist on the system.
3. **Create worker threads** to service completed I/O requests on the completion port using processor information in step 2.
4. Prepare a **listening socket** to listen for connections on port 5150.

5. Accept **inbound connections** using the **accept** function.
6. Create a data structure to represent **per-handle data** and **save the accepted socket handle** in the structure.
7. Associate the new socket handle returned from accept with the completion port by calling **CreateIoCompletionPort()**.
8. **Start processing I/O** on the accepted connection. Essentially, you want to post one or more asynchronous WSARecv() or WSASend() requests on the new socket using the overlapped I/O mechanism. When these I/O requests complete, **a worker thread services the I/O requests** and continues processing future I/O requests.
9. Repeat steps 5-8 until server terminates.

```

HANDLE CompletionPort;
WSADATA wsd;
SYSTEM_INFO SystemInfo;
SOCKADDR_IN InternetAddr;
SOCKET Listen;
int i;

typedef struct _PER_HANDLE_DATA
{
    SOCKET      Socket;
    SOCKADDR_STORAGE ClientAddr;
    // Other information useful to be associated with the handle
} PER_HANDLE_DATA, *LPPER_HANDLE_DATA;

// Load Winsock
StartWinsock(MAKEWORD(2, 2), &wsd);

// Step 1:
// Create an I/O completion port
CompletionPort = CreateIoCompletionPort(INVALID_HANDLE_VALUE, NULL, 0, 0);

// Step 2:
// Determine how many processors are on the system
GetSystemInfo(&SystemInfo);

```

```
// Step 3:
// Create worker threads based on the number of
// processors available on the system. For this
// simple case, we create one worker thread for each processor.
for (i = 0; i < SystemInfo.dwNumberOfProcessors; i++)
{
    HANDLE ThreadHandle;

    // Create a server worker thread, and pass the
    // completion port to the thread. NOTE: the
    // ServerWorkerThread procedure is not defined in this listing.
    ThreadHandle = CreateThread(NULL, 0, ServerWorkerThread, CompletionPort, 0, NULL;
    // Close the thread handle
    CloseHandle(ThreadHandle);
}
```

```
// Step 4:
// Create a listening socket
Listen = WSASocket(AF_INET, SOCK_STREAM, 0, NULL, 0, WSA_FLAG_OVERLAPPED);
```

```
InternetAddr.sin_family = AF_INET;
InternetAddr.sin_addr.s_addr = htonl(INADDR_ANY);
InternetAddr.sin_port = htons(5150);
bind(Listen, (PSOCKADDR)&InternetAddr, sizeof(InternetAddr));
```

```
// Prepare socket for listening
listen(Listen, 5);
```



```

while ( TRUE)
{
    PER_HANDLE_DATA *PerHandleData = NULL;
    SOCKADDR_IN saRemote;
    SOCKET Accept;
    int RemoteLen;

    // Step 5:
    // Accept connections and assign to the completion port
    RemoteLen = sizeof(saRemote);
    Accept = WSAAccept(Listen, (SOCKADDR *)&saRemote, &RemoteLen);

    // Step 6:
    // Create per-handle data information structure to associate with the socket
    PerHandleData = (LPPER_HANDLE_DATA) GlobalAlloc(GPTR, sizeof(PER_HANDLE_DATA));

    printf("Socket number %d connected\n", Accept);
    PerHandleData->Socket = Accept;
    memcpy(&PerHandleData->ClientAddr, &saRemote, RemoteLen);

    // Step 7:
    // Associate the accepted socket with the completion port
    CreateIoCompletionPort((HANDLE)Accept, CompletionPort, (DWORD)PerHandleData, 0);

    // Step 8:
    // Start processing I/O on the accepted socket.
    // Post one or more WSASend() or WSARecv() calls on the socket using overlapped I/O.
    WSARecv(...);
}

```

```
DWORD WINAPI ServerWorkerThread(LPVOID lpParam)
{
    // The requirements for the worker thread will be discussed later.
}
```

Completion Ports & Overlapped I/O

- ✓ After associating a socket handle with a completion port, you can begin processing I/O requests by posting overlapped send and receive requests on the socket handle.
- ✓ Winsock API calls such as **WSASend()** and **WSARecv()** return immediately when called.
- ✓ It is up to your application to retrieve the results of the calls at a later time through an **OVERLAPPED** structure.
- ✓ In the completion port model, this is accomplished by having **one or more worker threads** wait on the completion port using the **GetQueuedCompletionStatus()** function

```
BOOL GetQueuedCompletionStatus(  
    HANDLE CompletionPort,  
    LPDWORD lpNumberOfBytesTransferred,  
    PULONG_PTR lpCompletionKey,  
    LPOVERLAPPED * lpOverlapped,  
    DWORD dwMilliseconds  
);
```

- **CompletionPort**: the completion port to wait on.
- **lpNumberOfBytesTransferred**: receives the number of bytes transferred after a completed I/O operation, such as WSARecv() or WSASend().
- **lpCompletionKey**: returns per-handle data for the socket that was originally passed into the CreateIoCompletionPort() function.
 - recommend saving the socket handle in this key.
- **lpOverlapped**: receives the WSAOVERLAPPED structure of the completed I/O operation
- **dwMilliseconds**: the number of milliseconds that the caller is willing to wait for a completion packet to appear on the completion port.

Per-handle Data and Per-I/O Operation Data

- ✓ The **IpCompletionKey** parameter contains what we call **per-handle data** because the data is related to a socket handle when a socket is first associated with the completion port.
- ✓ The **IpOverlapped** parameter contains an OVERLAPPED structure followed by what we call **per-I/O operation data**, which is anything that your worker thread will need to know when processing a completion packet.

```
typedef struct
{
    OVERLAPPED Overlapped;
    char      Buffer[DATA_BUFSIZE];
    int       BufferLen;
    int       OperationType;
} PER_IO_DATA;
```

```
PER_IO_OPERATION_DATA PerloData;
```

```
WSABUF wbuf;
```

```
DWORD Bytes, Flags;
```

```
// Initialize wbuf ...
```

```
WSARecv(socket, &wbuf, 1, &Bytes, &Flags, &(PerloData.Overlapped), NULL);
```

```
//Later in the worker thread, GetQueuedCompletionStatus() returns with an  
//overlapped structure and completion key. To retrieve the per-I/O data the macro  
//CONTAINING_RECORD should be used. For example,
```

```
PER_IO_DATA *PerloData=NULL;
```

```
OVERLAPPED *lpOverlapped=NULL;
```

```
ret = GetQueuedCompletionStatus(  
    CompPortHandle,  
    &Transferred,  
    (PULONG_PTR)&CompletionKey,  
    &lpOverlapped, INFINITE);
```

```
// Check for successful return
```

```
PerloData = CONTAINING_RECORD(lpOverlapped, PER_IO_DATA, Overlapped);
```

Practice

✓ chapter05 → iocpServer project

I/O Model Consideration

- ✓ When you are developing a **client** application that manages one or more sockets, we recommend using overlapped I/O or `WSAEventSelect()` over the other I/O models for performance reasons.
- ✓ When you are developing a **server** that processes several sockets at a given time, we recommend using overlapped I/O over the other I/O models for performance reasons.

References

- ✓ <http://www.madwizard.org/programming/tutorials/netcpp/5>

MY **BRIGHT** FUTURE

DSU Dongseo University
동서대학교